

Apache Cassandra

A NoSQL Database

What is Cassandra?

- The Apache Cassandra is a NoSQL database which provides scalability and high availability without compromising performance.
- Linear scalability and proven fault-tolerance on commodity hardware make it the perfect platform for mission-critical data.
- Cassandra's support for replicating across multiple data centers is best-in-class.

A Brief History

- Initially developed at Facebook for inbox search.
- Open sourced by Facebook in 2008.
- Accepted into Apache Incubator in 2009.
- Apache top-level project since February 2010.
- Now maintained by the Apache Software Foundation.

Features

- **Fault Tolerant**
 - Data is automatically replicated to multiple nodes for fault-tolerance.
 - Failed nodes can be replaced with no downtime.
- **Performant**
 - Cassandra consistently outperforms popular NoSQL alternatives in benchmarks.
- **Decentralized**
 - There are no single points of failure. There are no network bottlenecks. Every node in the cluster is identical.

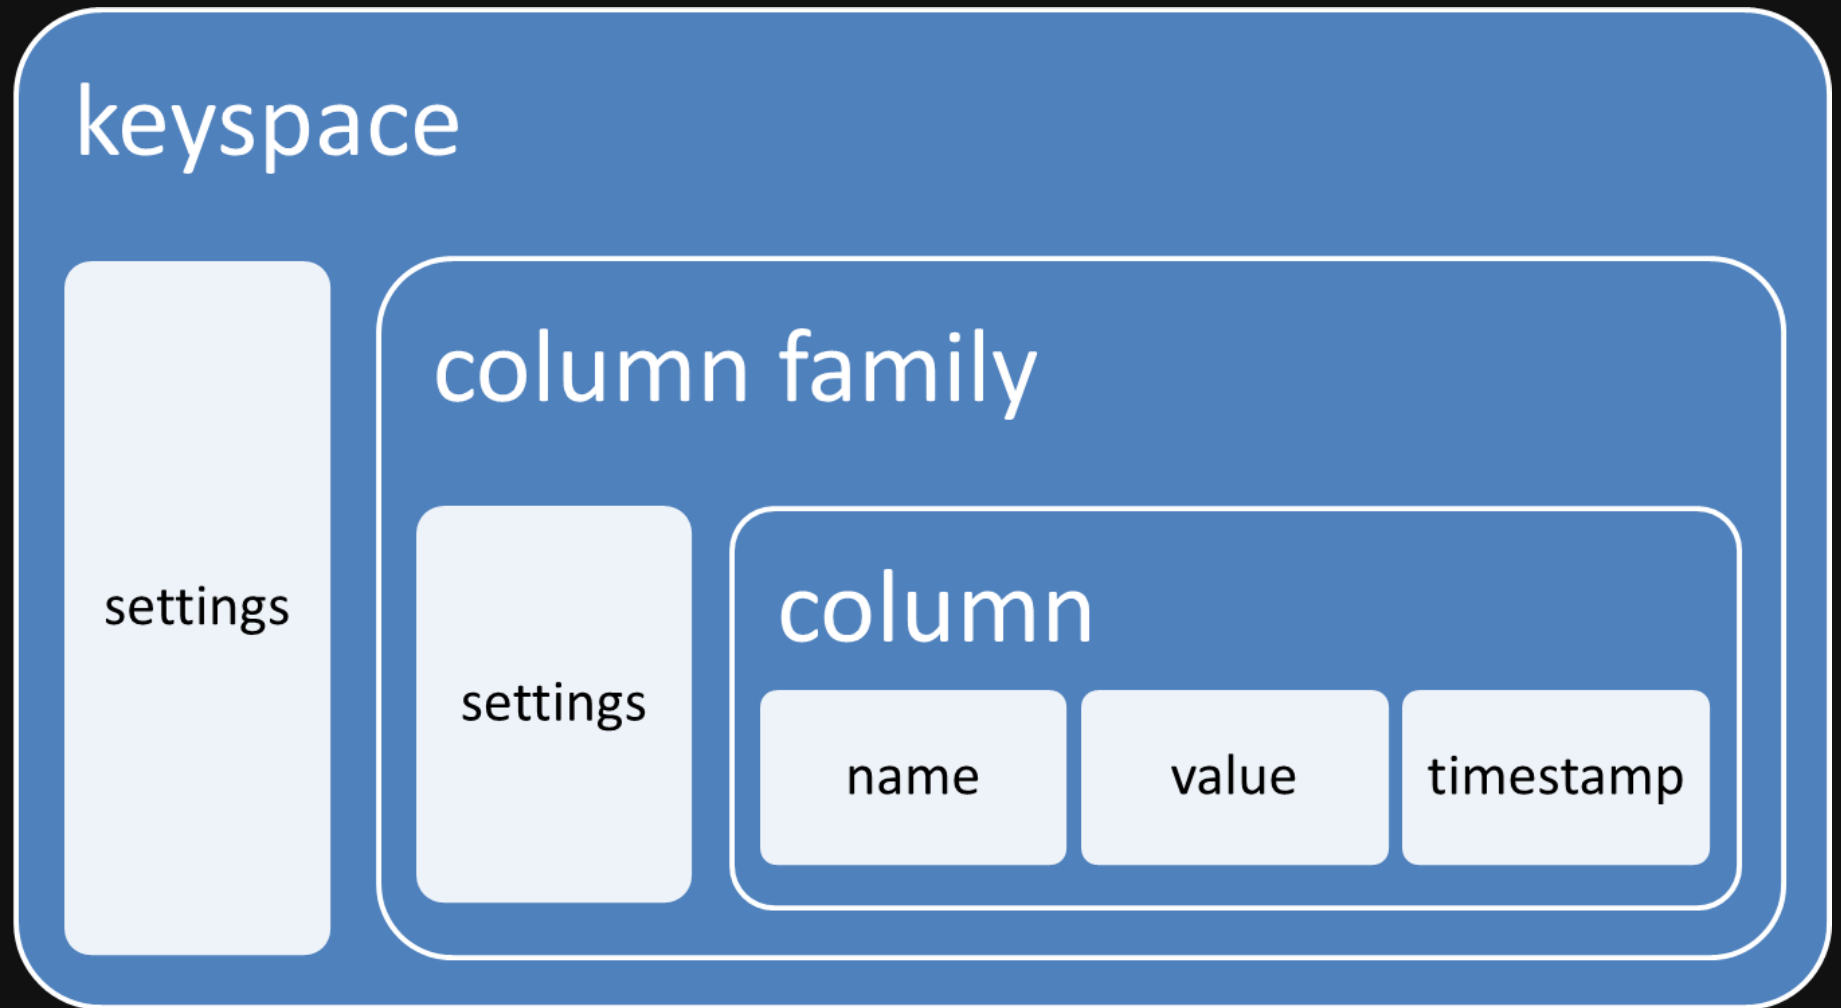
- **Scalable**

- Some of the largest production deployments include:
- Apple's, with over 75,000 nodes storing over 10 PB of data,
- Netflix (2,500 nodes, 420 TB, over 1 trillion requests per day),
- eBay (over 100 nodes, 250 TB).

- **Durable**

- Cassandra is suitable for applications that can't afford to lose data, even when an entire data center goes down.

Data Model



Data Model

- Keyspace
 - A keyspace defines a number of options that applies to all the tables it contains.
 - Most prominent is the replication strategy.
- Table
 - A multi-dimensional map indexed by key which stores the rows.
- Column
 - Basic data structure that stores the data elements.
 - Each column contains name, value and timestamp.

Cassandra Query Language (CQL)

- Offers a syntax similar to SQL.
- Offers a data model similar to SQL so that data is stored in rows of columns.
- Provides:
 - Data Definition Statements
 - Data Manipulation Statements
- Using the command **cqlsh**.

Type	Description
blob	Arbitrary Bytes
boolean	<i>True</i> or <i>False</i> values
date	Date (without time value)
timestamp	Date (with time value)
double	64-bit Floating Point
float	32-bit Floating Point
int	32-bit Signed Integer
text	UTF8 Encoded String
varchar	UTF8 Encoded String

Type	Description
map<T, T>	A sorted set of key-value pairs
set<T>	A sorted collection of unique values
list<T>	A collection of non-unique values
tuple<T, T>	A pair of values (a, b)

Creating a Keyspace

```
CREATE KEYSPACE Emp  
  WITH replication = {'class': 'SimpleStrategy', 'replication_factor': 1};
```

- Creates a keyspace named **Emp**
- replication_factor determines number of copies.

Switch to created keyspace using:

```
USE Emp;
```

Creating a Table

```
CREATE TABLE [table name] (  
    [column name] [column type] [PRIMARY KEY]  
) WITH [table options];
```

Example

```
CREATE TABLE Employee (  
    id int PRIMARY KEY,  
    name text,  
    city text,  
    salary int);
```

Insert Statement

```
INSERT INTO [table name] ([column names])  
VALUES ([values]);
```

Example

```
INSERT INTO Employee (id, name, city, salary)  
VALUES (1, 'Name 1', 'Pune', 25000);
```

- If the row with given primary key does not exist, it will be created.
- Otherwise, the existing row will be updated.

Select Statement

```
SELECT [column names] FROM [table]  
      WHERE [where clause]  
      ORDER BY [column name [ASC | DESC]]  
      LIMIT [limit];
```

Example

```
SELECT id, name, city, salary FROM Employee;
```

Update Statement

```
UPDATE [table name]  
    SET [column name] = [value]  
    WHERE [where clause];
```

Example

```
UPDATE Employee SET salary = 50000 WHERE id = 4;
```

Delete Statement

```
DELETE FROM [table name] WHERE [where clause];
```

Example

```
DELETE FROM Employee WHERE id = 1;
```

- If column name is specified after DELETE keyword, only that column is removed.
- Otherwise the entire row is removed.

Aggregate Functions

- Aggregate functions work on a set of rows to deliver an aggregated result.
- Following built in aggregation functions are available:
 - **count()**: Count rows returned by the query.
 - **max()** and **min()**: Compute maximum and minimum value returned by a query.
 - **sum()**: Sum up all values returned by a query.
 - **avg()**: Average of all values returned by a query.

Aggregation Examples

Count:

```
SELECT count(*) FROM Employee;
```

Min and Max:

```
SELECT max(salary) FROM Employee;  
SELECT min(salary) FROM Employee;
```

Sum:

```
SELECT sum(salary) FROM Employee WHERE city = 'Pune';
```

Avg:

```
SELECT avg(salary) FROM Employee WHERE city = 'London';
```

Indexing

- Indexed using primary key by default.
- Secondary indexes can be created as required with the syntax:

```
CREATE INDEX [index name] on [table name] ([column]);
```

```
CREATE INDEX nameIndex on Employee (name);
```

- Indexes can be dropped by using:

```
DROP INDEX [index name];
```

```
DROP INDEX nameIndex;
```